Production Deployment

MLOps: Training Operationalization

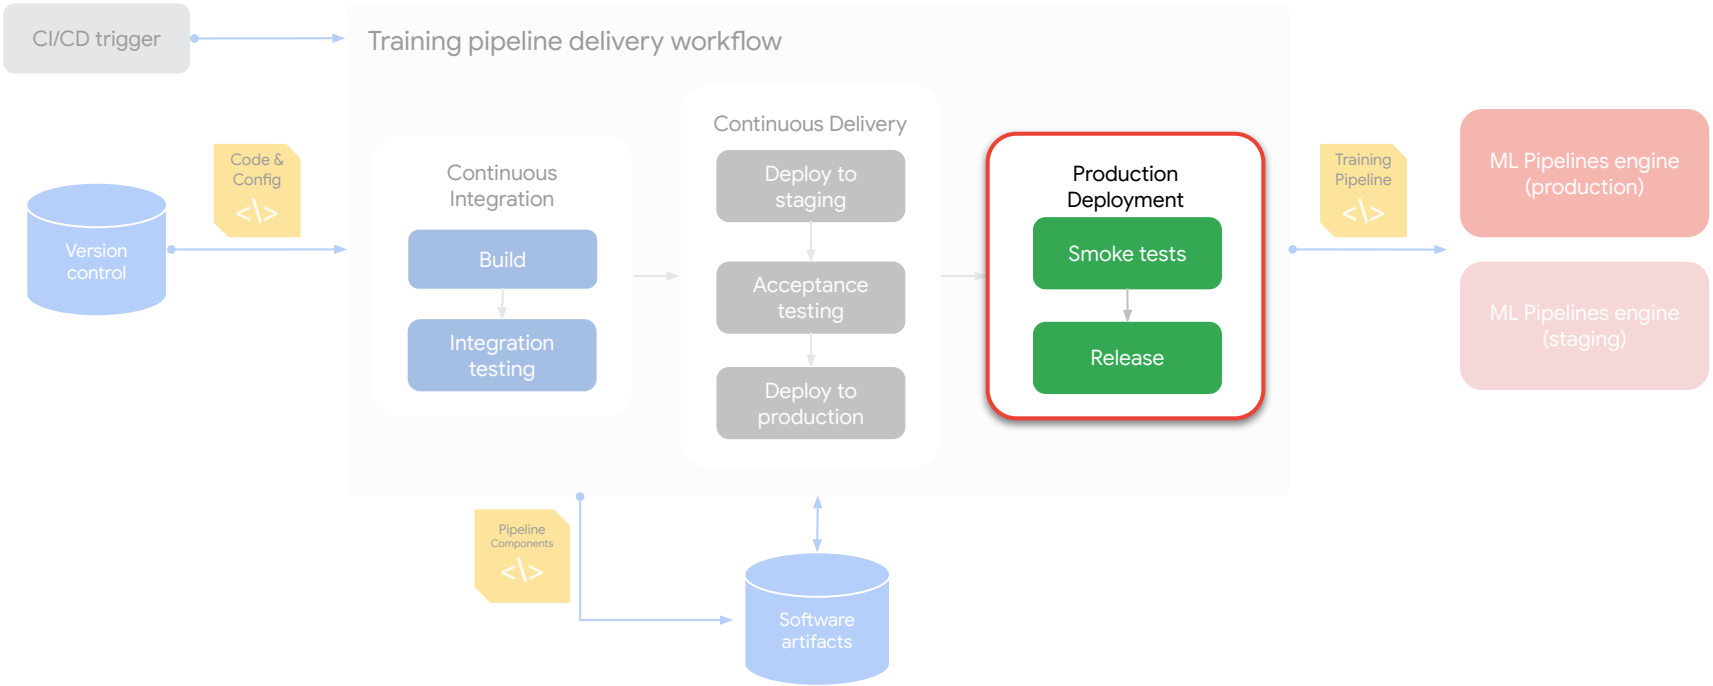

The MLOps **Personas**

ML Engineer

ML Researcher

Data Scientist

Data Engineer

Software Engineer

DevOps

Business Analyst

Things go wrong even ***after*** testing?

Things go wrong even ***after*** testing?

1. **Assumed** testing environments is the **same** as the production environment

Things go wrong even ***after*** testing?

1. **Assumed** testing environments is the **same** as the production environment

2. **Hard to** make the test environment **match** the production environment

Things go wrong even ***after*** testing?

1. **Assumed** testing environments is the **same** as the production environment
2. **Hard to** make the test environment **match** the production environment
3. **Impossible to test all** the different **conditions** and scenarios in practice

Things go wrong even ***after*** testing?

1. **Assumed** testing environments is the **same** as the production environment
2. **Hard to** make the test environment **match** the production environment
3. **Impossible to test all** the different **conditions** and scenarios in practice

⇒ *"Test" live in production environment with real-traffic*

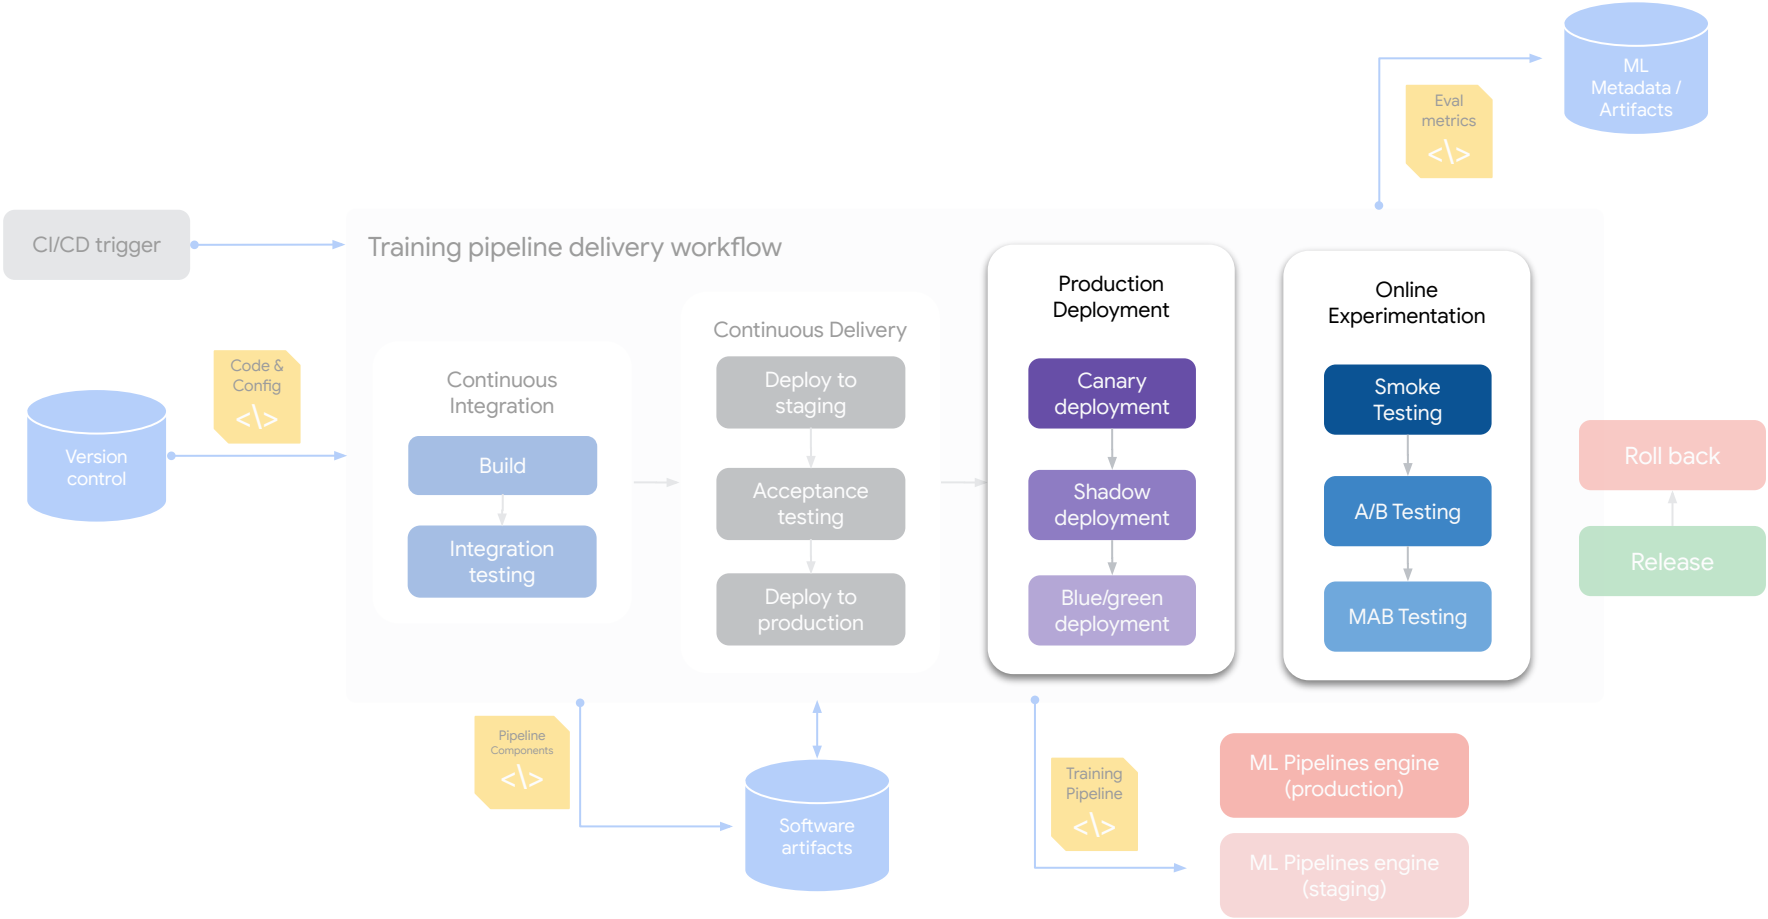

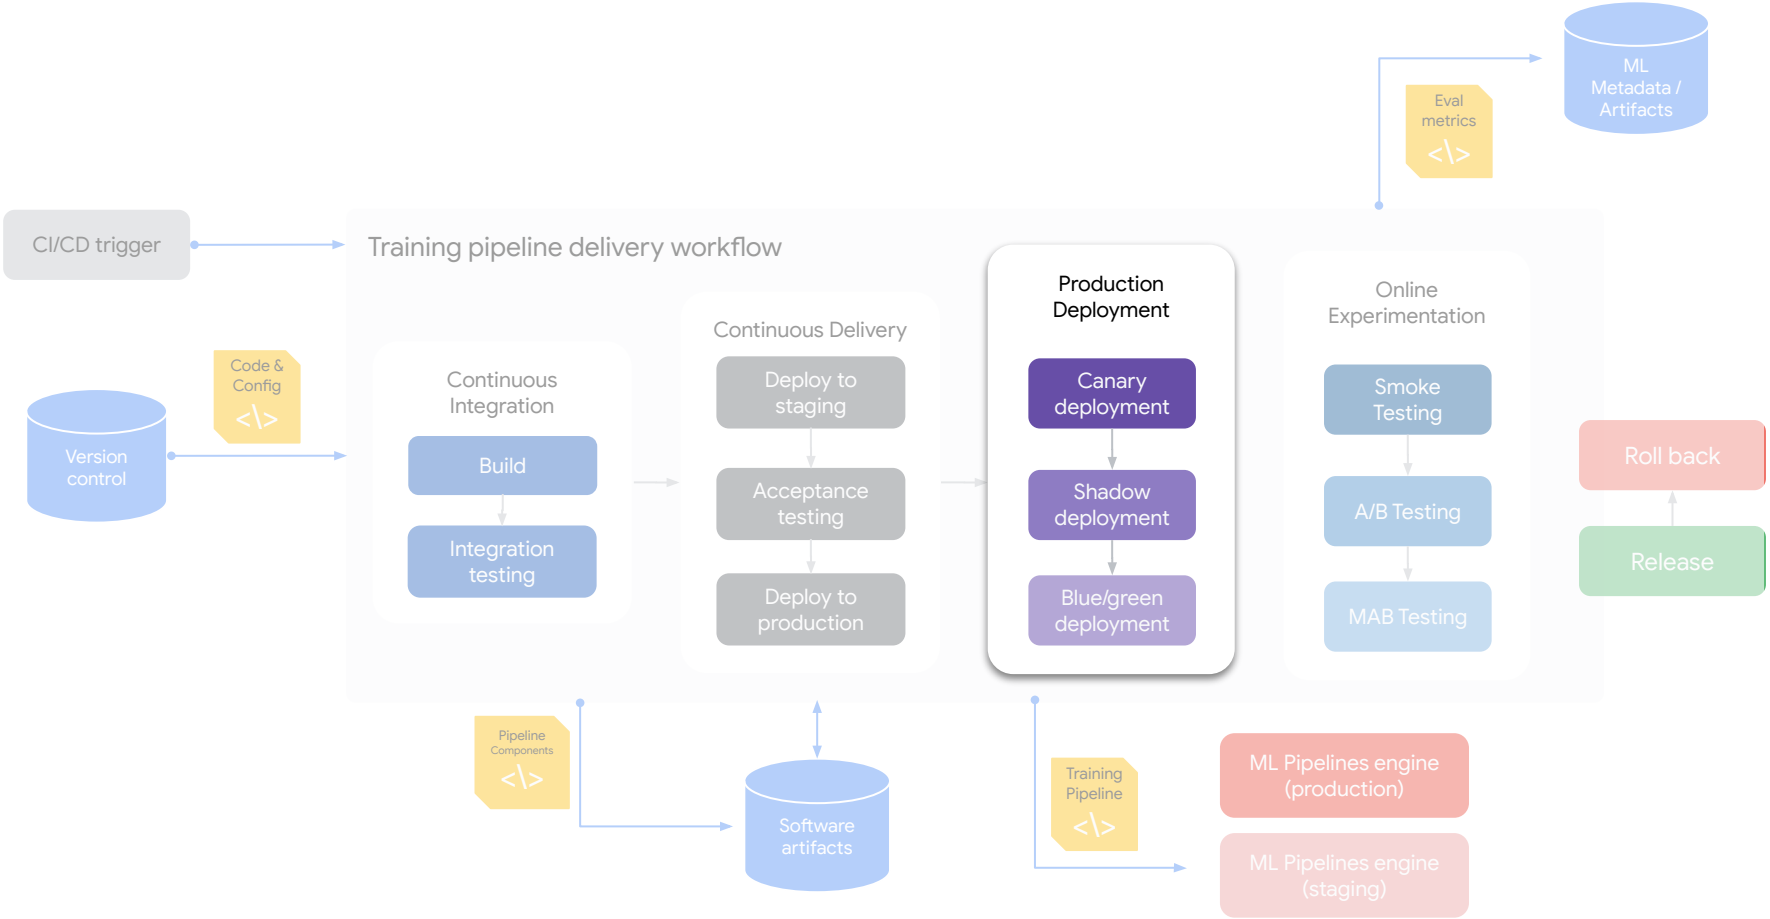

Production Deployment Strategies

1. Canary Deployment
2. Shadow Deployment
3. Blue/Green Development

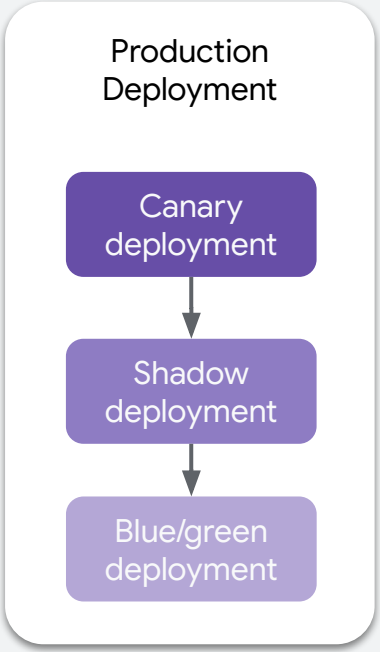

Canary Deployment

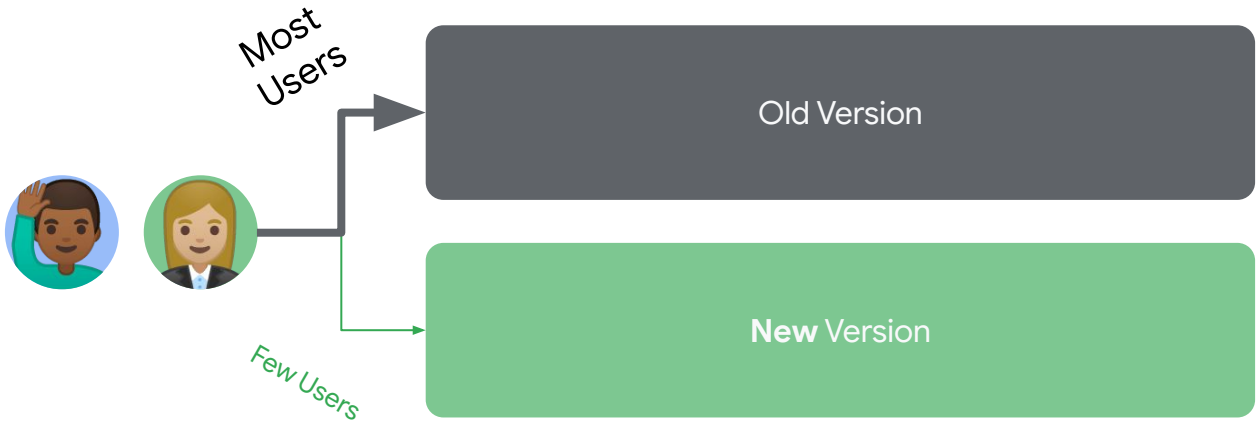

Canary Deployment

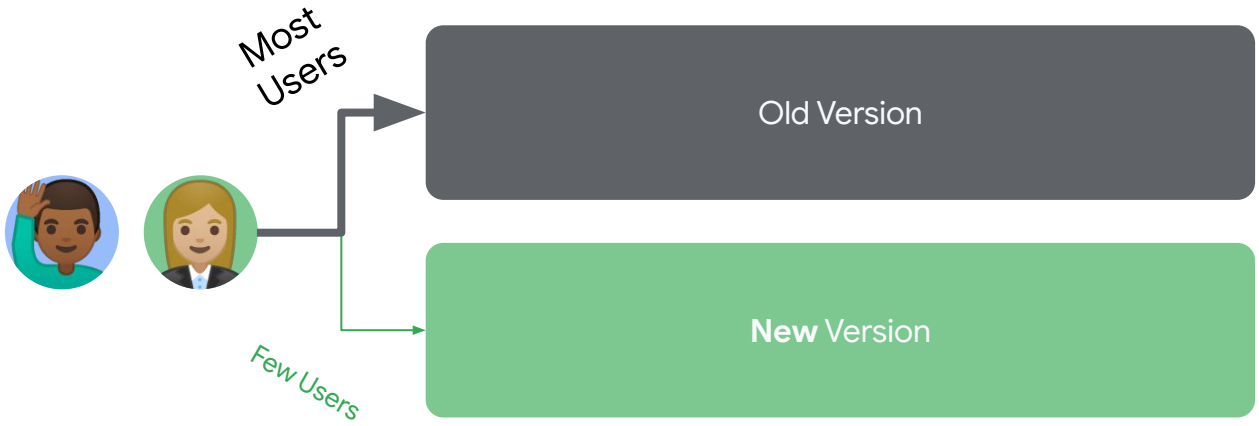

- Few affected by problems

Canary Deployment

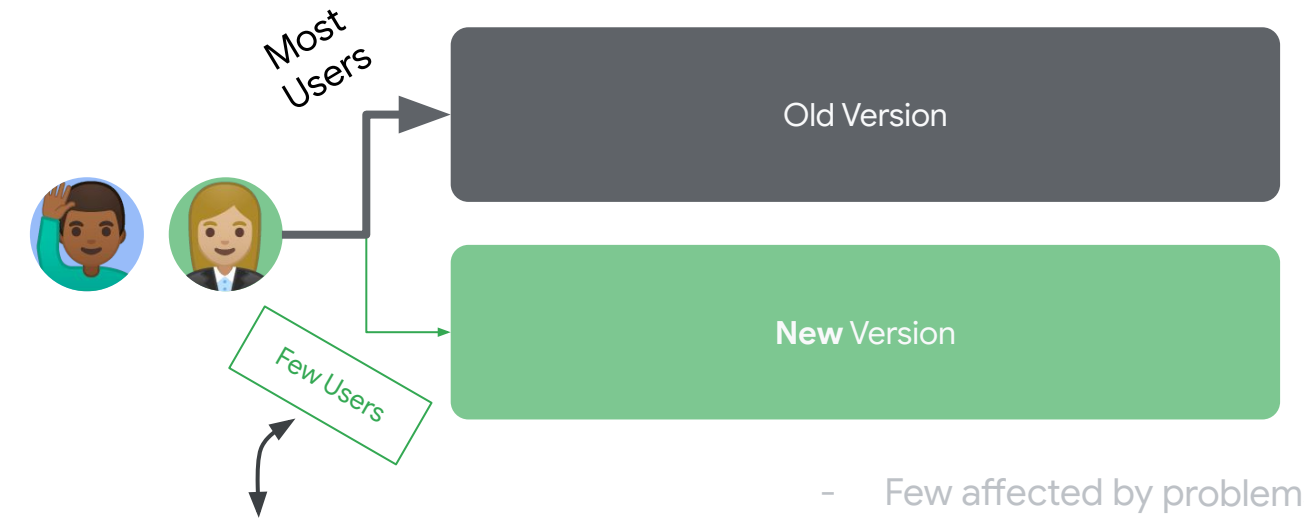

- Few affected by problems

Shadow Deployment

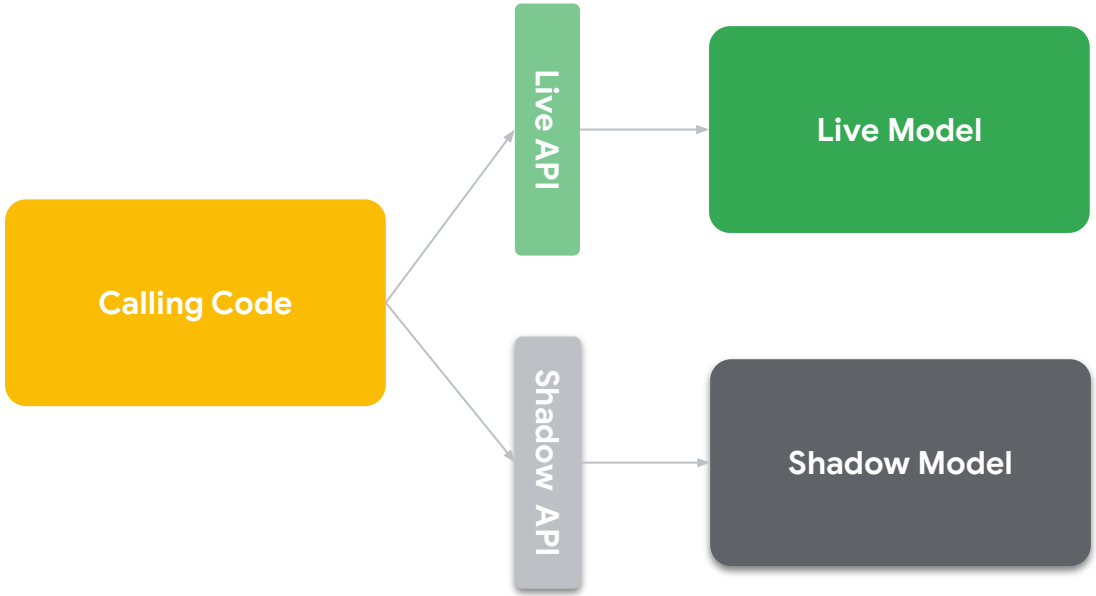

Shadow Deployment

+ **No impact to production** traffic (customer)
+ **Check stability** (requirements)

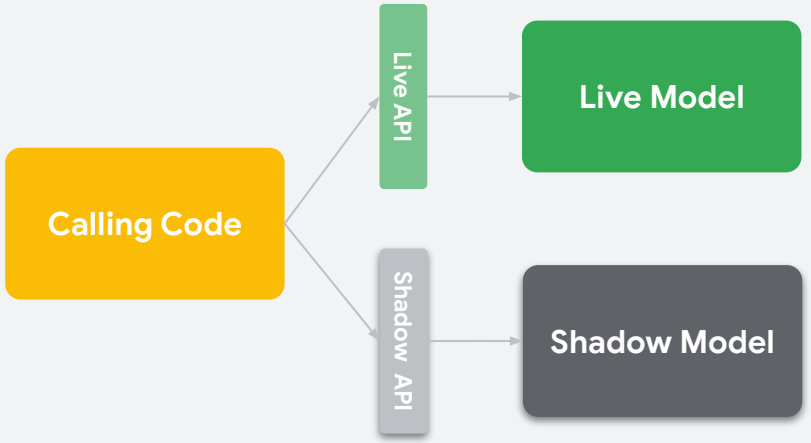

Shadow Deployment

+ **No impact to production** traffic (customer)
+ **Check stability** (requirements)

- Increased **cost** (parallel)
- Complex **update-in-place**
- Need to carefully design **performance testing**

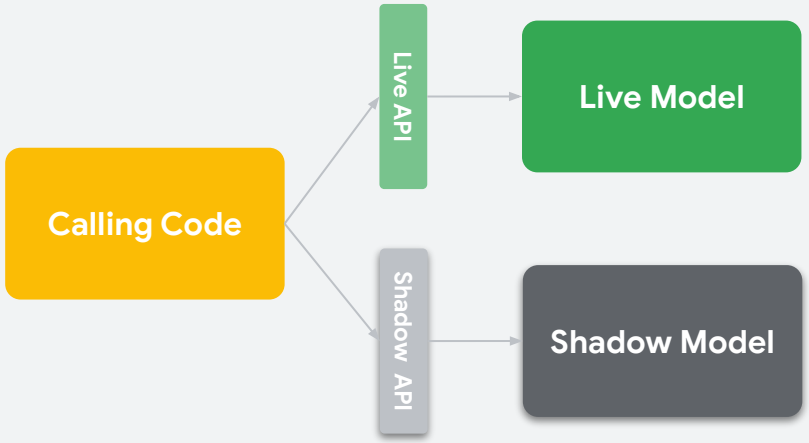

Blue/Green Deployment

In its purest form, blue-green asks us to **duplicate** every resource our application leverages.

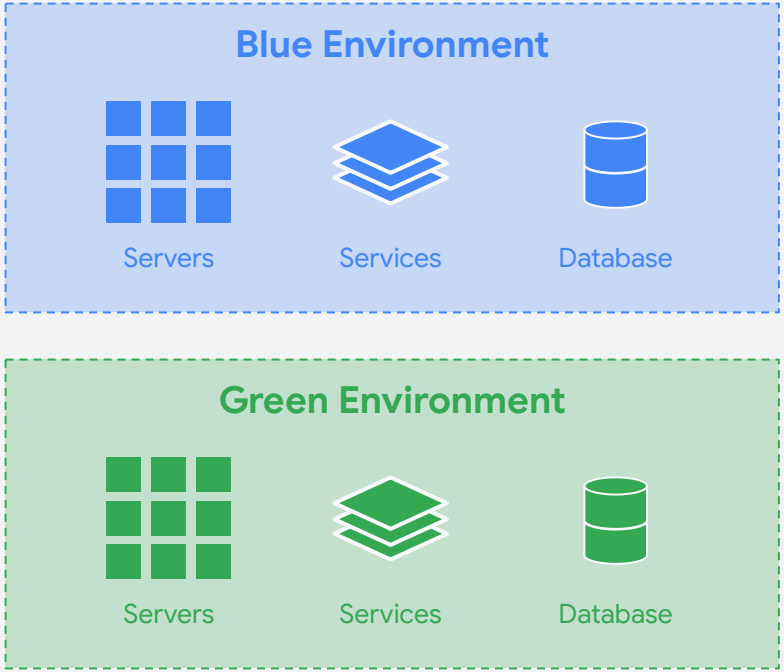

Blue/Green Deployment

In its purest form, blue-green asks us to **duplicate** every resource our application leverages.

But, we often see blue-green deployments with **shared** components.

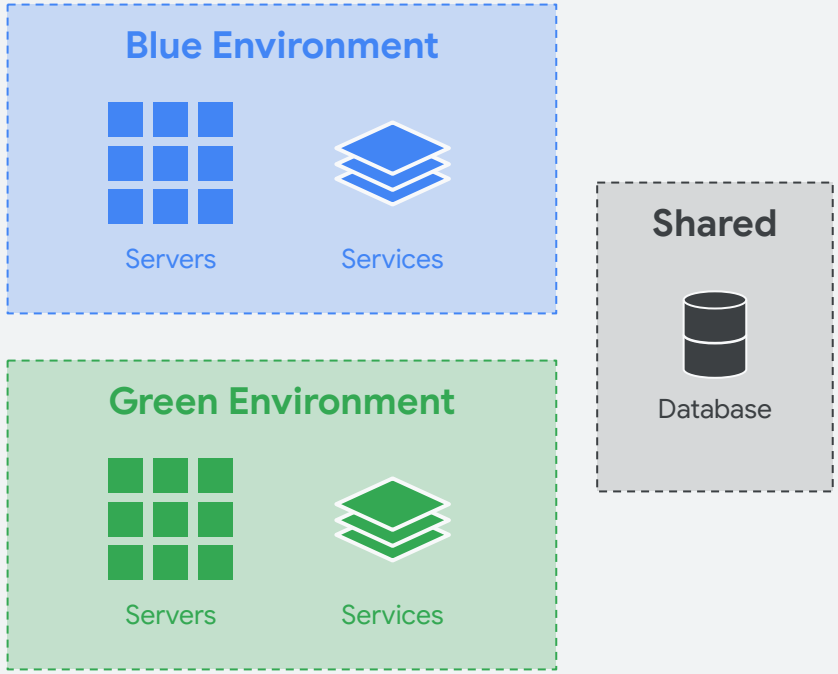

Blue/Green Deployment

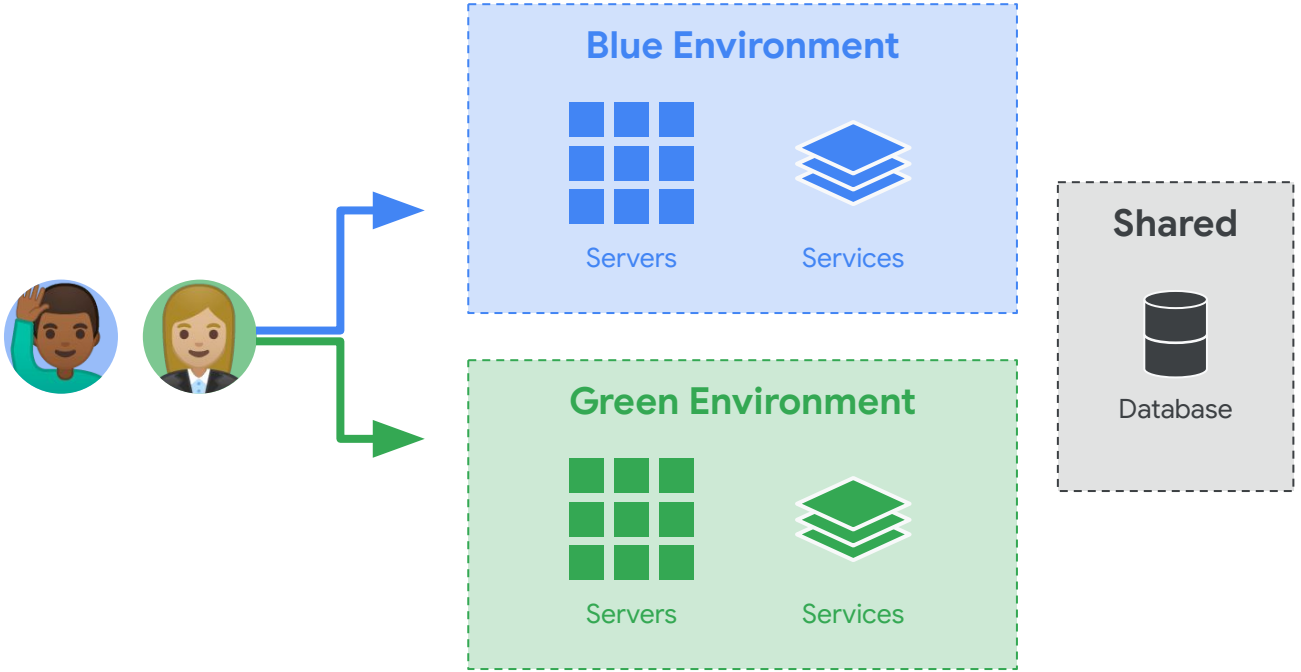

Blue/Green Deployment:
Alternate Producers

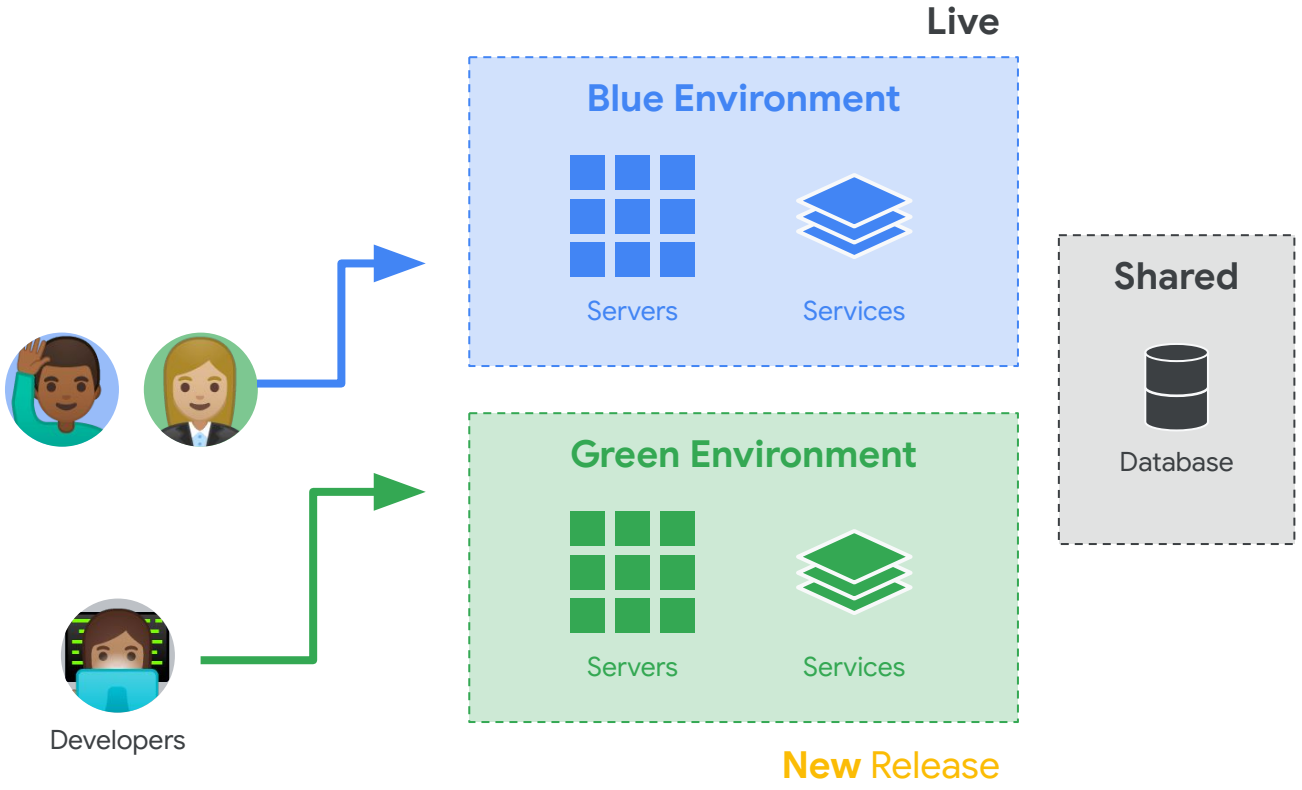

Blue/Green Deployment:
Alternate Producers

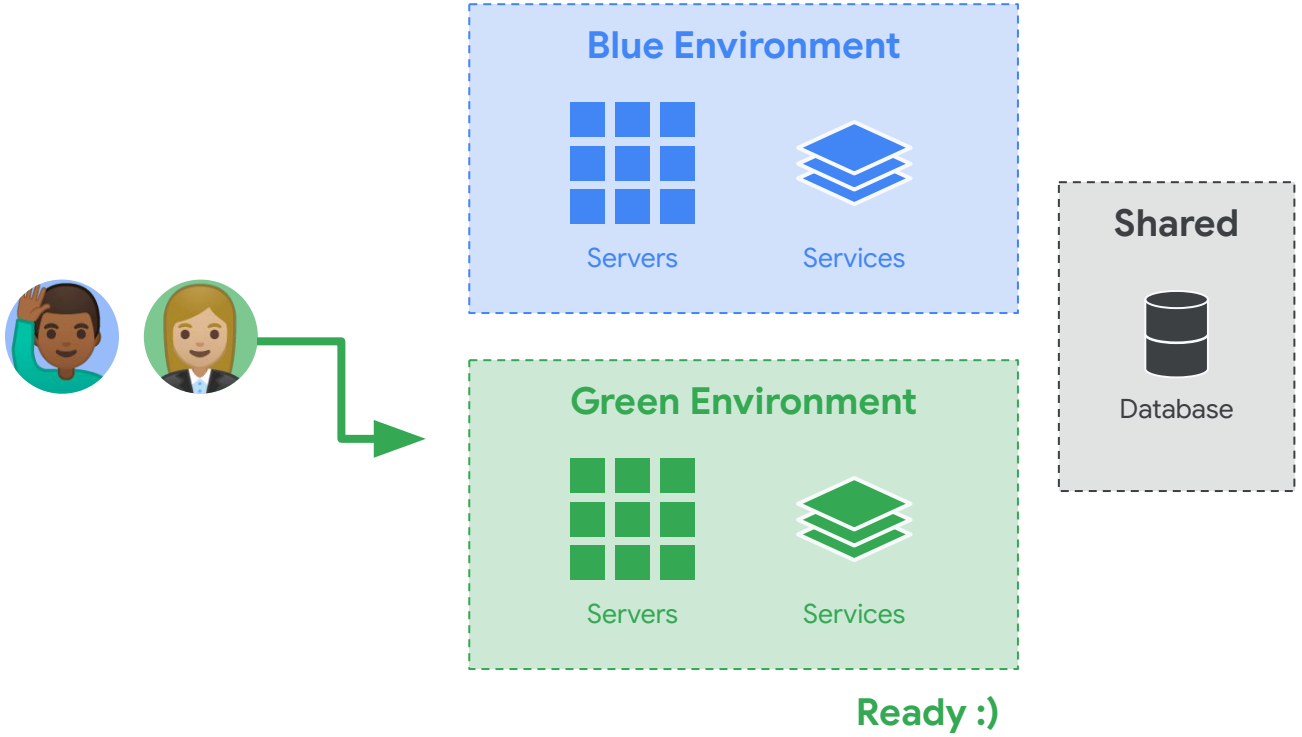

Blue/Green Deployment

+ No system downtime
+ Accurate testing
+ Reliability testing of deployments

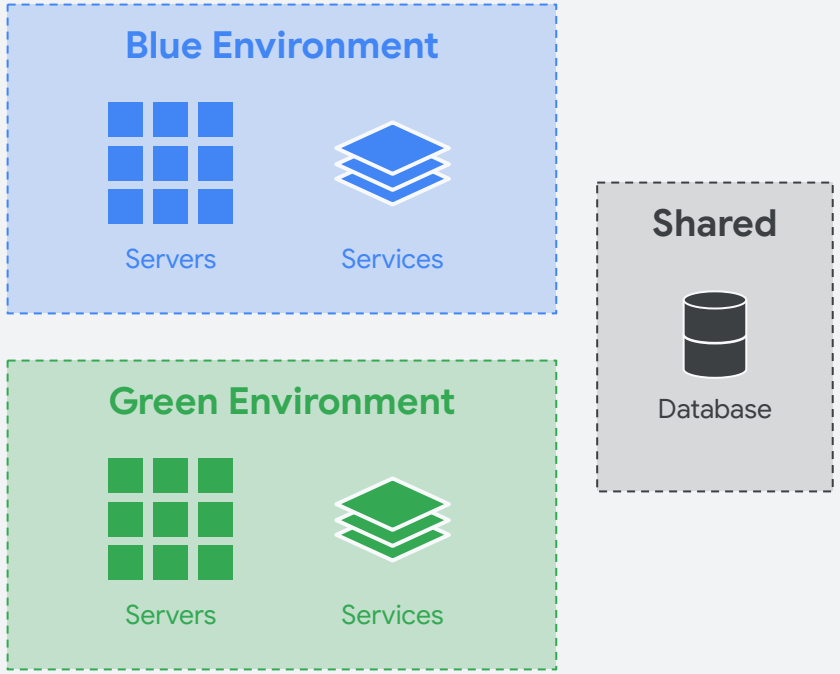

Blue/Green Deployment

For **TinyML**:

+ **When we can't make continuous updates**

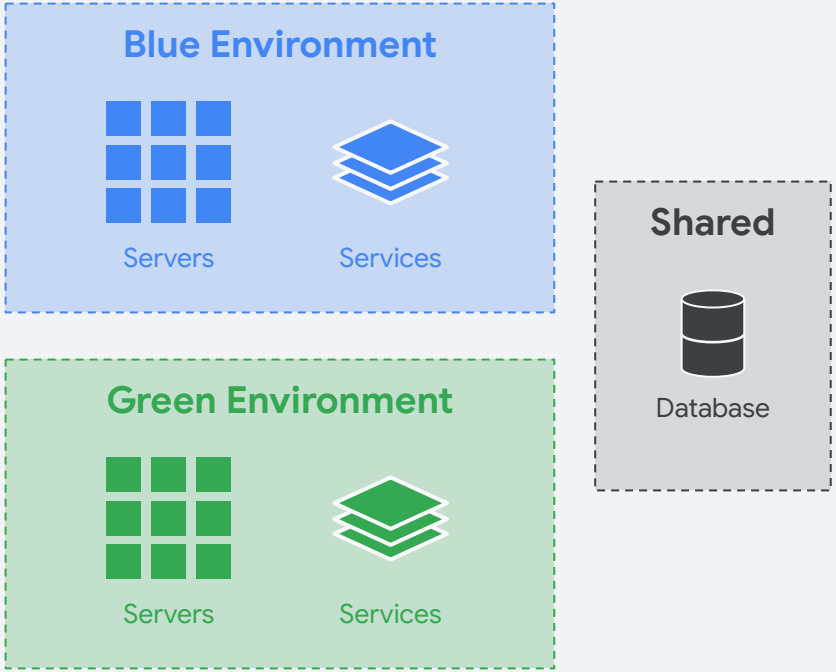

Blue/Green Deployment

For **TinyML**:

+ When we can't make continuous updates
+ **When we can't support two identical environments**

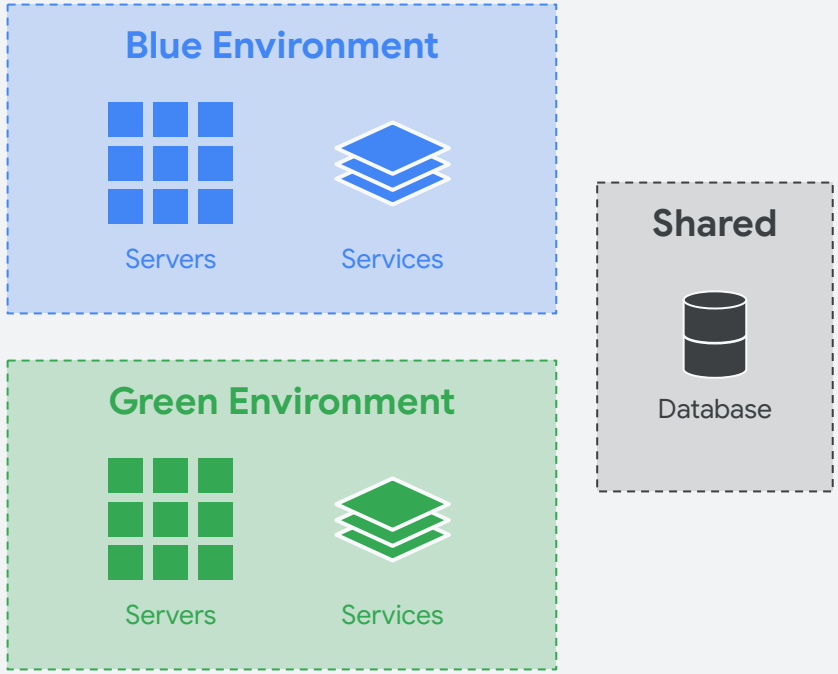

Blue/Green Deployment

For **TinyML**:

+ When we can't make continuous updates
+ When we can't support two identical environments
+ **When infrastructure does not allow a router to direct users**